

COMP9334

Project, Session 1 2017

Getting the best response time out of your power bill

Elliot Griffiths

z3332776

Generating Probability Distributions

To generate the inter-arrival times of the requests arriving at the high-speed router we take the product of two random numbers, a_{1k} and a_{2k} . Such that the inter-arrival time is $a_k = a_{1k}a_{2k}$.

To generate the exponentially distributed values a_{1k} we use the inverse transform method. We generate a random number u , which is uniformly distributed in $(0,1)$ and compute the number $F^{-1}(u)$ where $F(x)$ is the Cumulative Density Function (CDF)

We know the CDF for an exponential distribution is $F(x) = 1 - e^{-\lambda x}$
 Thus, $u = 1 - e^{-\lambda x}$ and therefore our generated number is $x = -\frac{\log(1-u)}{\lambda}$
 (where $\lambda = 7.2$)

To generate the uniformly distributed values a_{2k} in the interval $[0.75, 1.17]$ we simply generate a number between 0 and $1.17 - 0.75$ then we add that number to 0.75 giving us the required range.

Finally generating the service time t is slightly more involved. We are given the probability density function $g(t)$.

Where Probability Density Function of service time is:

$$g(t) = \begin{cases} 0 & \text{for } \alpha_1 \leq t \\ \frac{\gamma}{t^\beta} & \text{for } \alpha_1 \leq t \leq \alpha_2 \\ 0 & \text{for } t \geq \alpha_2 \end{cases}$$

$$\text{Where, } \gamma = \frac{1-\beta}{\alpha_2^{1-\beta} - \alpha_1^{1-\beta}}$$

For generating random numbers with this probability distribution, we first need to find the Cumulative Density Function $G(t)$. To do this we integrate over the curve $\frac{\gamma}{t^\beta}$ from α_1 to t

This gives the $CDF(t) = G(t)$

$$G(t) = \begin{cases} 0 & \text{for } t < \alpha_1 \\ \frac{t^{1-\beta} - \alpha_1^{1-\beta}}{(\alpha_2^{1-\beta} - \alpha_1^{1-\beta})} & \text{for } \alpha_1 \leq t \leq \alpha_2 \\ 1 & \text{for } t \geq \alpha_2 \end{cases}$$

We can see that if $t = \alpha_2$ then the cumulative probability of getting a value of 0.98 or less is 100% thus it is correct.

$$u = \frac{t^{1-\beta} - \alpha_1^{1-\beta}}{\alpha_2^{1-\beta} - \alpha_1^{1-\beta}}$$

Now we need to arrange the function so we can find the 't-intercepts' from given 'u' values

Therefore our values for t given u between 0 and 1 are

$$t = \left(u * (\alpha_2^{1-\beta} - \alpha_1^{1-\beta}) + \alpha_1^{1-\beta} \right)^{\frac{1}{1-\beta}}$$

Simulating the PS Server

The simulation is written in perl and stored in file `ps_server_sim.pl`. To run the simulation two arguments must be given after the program name. The first is the simulation time and the second is the number of servers. For example:

```
Elliot@MacBook-Pro:Project Elliot$ ./ps_server_sim.pl 10000 10
```

```
7495 jobs completed by 1 server so we have 74950 jobs completed by
10 servers,
in 4064.88435799293 time therefore mean response time is
0.542346145162499
```

Verifying the Simulation

To verify the code I have wrote an alternative versions of the simulation with the random number generation removed and replaced with two lists containing the Arrival time of the request and the service time of each request. These values are popped from the lists when required. For simplicity the frequency has been set to 1GHz

CODE: ps_server_sim_ver.pl

SERVERS: 1 at 1GHz

Job Number	Request Arrival Time	Service Time
1	1	2.1
2	2	3.3
3	3	1.1
4	5	0.5

Expected Mean Response Time: 3.6

Mean Response Time Simulation:

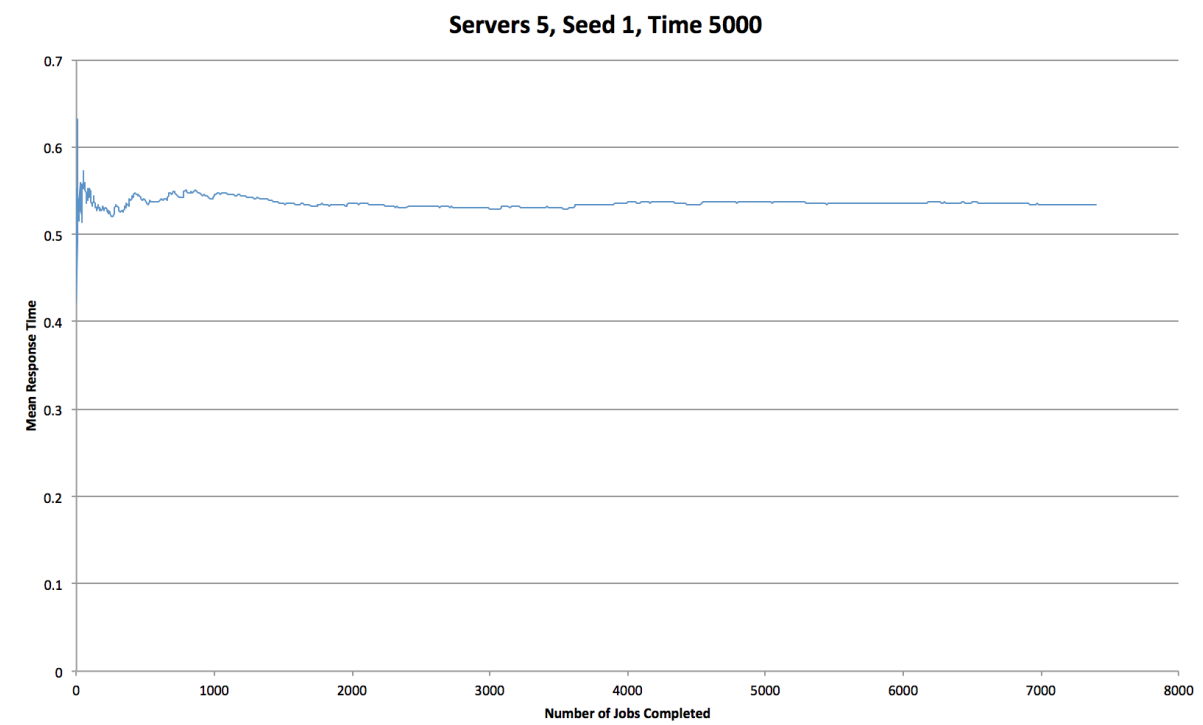
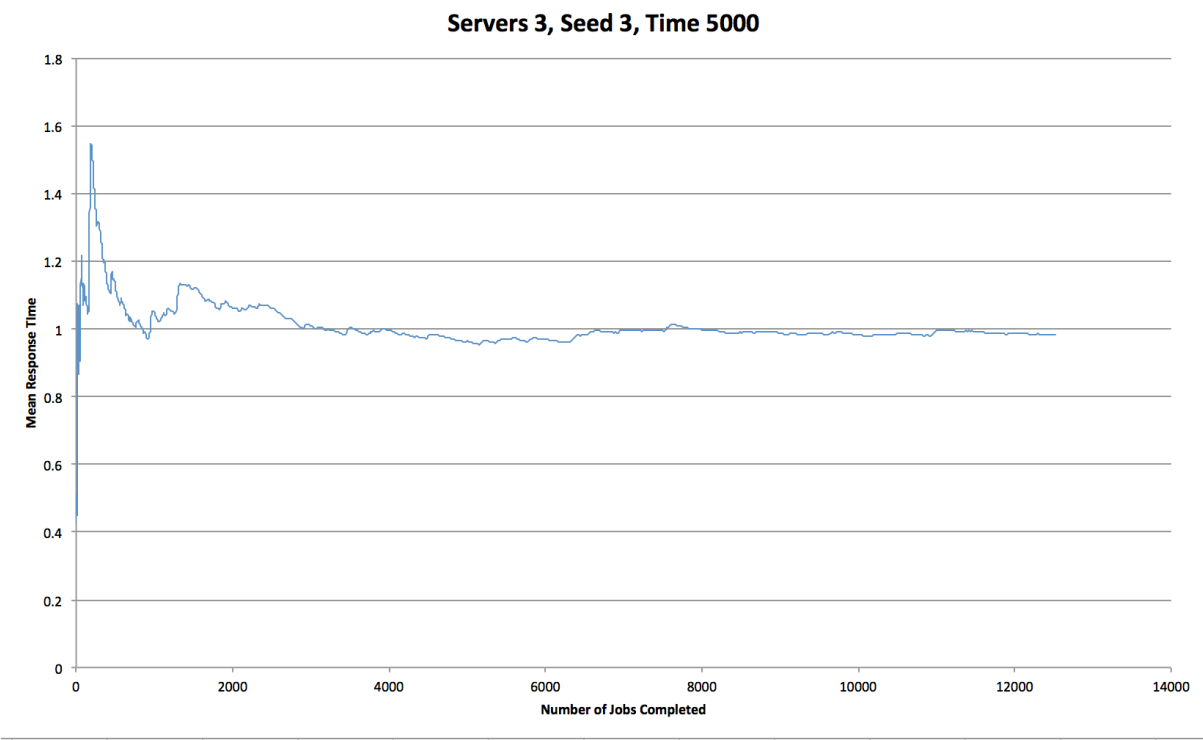
```
Elliot's-MacBook-Pro:Validating Elliot$ ./ps_server_sim_ver.pl
```


```
4 jobs completed by 1 server(s) at 1GHz frequency in 14.4 time  
therefore mean response time is 3.6
```

Choosing Simulation Parameters

SIM TIME

First we want to make sure we have simulated enough jobs such that the mean response time has reached steady state behaviour and does not show any transient behaviour. We do this by visually inspecting the mean response time as the number of jobs completed increases.







We can see from inspection of the above graphs that the response time leaves the transient section and reaches steady state after around 4000 jobs completed.

To ensure our simulated response times have left the transient section we will start recording the response times after 5000 time units have past.

We now need to decide how long to run the simulation. i.e From 5000 time units to X time units.





From inspection we can see that after removal of the first 5000 time units the mean response time returns to a steady state after a following 10000 time units. Thus we will simulate for 0-15000 time units and remove the first results from 0 - 5000 time units.

INDEPENDENT REPLICATIONS

Now we need to find an appropriate number of independent replications to achieve a confidence interval of 95% that has a small enough range to avoid overlap with the confidence intervals simulations.

First we will try with $n = 30$ independent replications. If we need to tighten the confidence interval we can either increase n or simulate for longer time to reduce the standard deviation.

3 Servers

Seed	Mean Response Time	Sample Mean
1	0.968082436	1.009708383
2	1.01165843	
3	1.003298418	Sample SD
4	0.978125311	0.026795794
5	0.977540504	
6	1.025093734	alpha
7	1.047546367	0.05
8	1.035121274	
9	1.035590128	scale factor (t 29, 0.975)
10	1.005032247	2.04523
11	0.975570673	
12	1.022300154	+/- Scale factor * (S / \sqrt{n})
13	1.00178242	0.010005716
14	0.971986812	
15	1.023872489	
16	0.998126317	Lower bound
17	1.001315833	0.999702667
18	0.982632928	
19	1.013761194	Upper bound
20	1.018326441	1.019714098
21	1.010792162	
22	0.997688981	
23	1.070610585	
24	0.97995469	
25	1.033074155	
26	1.046638712	
27	1.060239365	
28	1.009126017	
29	0.982738248	
30	1.00362445	

4 Servers

Seed	Mean Response Time	Sample Mean
1	0.63464628	0.630695443
2	0.61785193	
3	0.632199762	Sample SD
4	0.631507602	0.006631602
5	0.633128034	
6	0.636228298	alpha
7	0.63478278	0.05

8	0.623928146	
9	0.638921156	scale factor (t 29, 0.975)
10	0.629290194	2.04523
11	0.626338276	
12	0.635655258	+/- Scale factor * (S / \sqrt{n})
13	0.627965373	0.002476281
14	0.633041872	
15	0.64273823	
16	0.631995816	Lower bound
17	0.627915688	0.628219162
18	0.635698368	
19	0.634197564	Upper bound
20	0.620890478	0.633171724
21	0.629271311	
22	0.611771915	
23	0.637602432	
24	0.625897801	
25	0.625393143	
26	0.62595185	
27	0.630068343	
28	0.634728134	
29	0.631629221	
30	0.639628035	

5 Servers		
Seed	Mean Response Time	Sample Mean
1	0.552066987	0.543144946
2	0.545688779	
3	0.541961551	Sample SD
4	0.547222257	0.003528652
5	0.545623846	
6	0.538960804	alpha
7	0.535176456	0.05
8	0.543329568	
9	0.542350073	scale factor (t 29, 0.975)
10	0.539715454	2.04523
11	0.543365202	
12	0.542884593	+/- Scale factor * (S / \sqrt{n})
13	0.547354775	0.001317621
14	0.542595902	
15	0.544027293	
16	0.543604611	Lower bound
17	0.544115126	0.541827326
18	0.542584239	
19	0.540711322	Upper bound
20	0.534807522	0.544462567

21	0.543894231
22	0.545053141
23	0.540259757
24	0.538289794
25	0.546829823
26	0.544219026
27	0.543302512
28	0.544165087
29	0.544070339
30	0.546118319

6 Servers

Seed	Mean Response Time	Sample Mean
1	0.519493955	0.517849328
2	0.512292896	
3	0.515657311	Sample SD
4	0.516548809	0.002847537
5	0.518659901	
6	0.520016773	alpha
7	0.519386681	0.05
8	0.521785967	
9	0.519000391	scale factor (t 29, 0.975)
10	0.51306208	2.04523
11	0.522765528	
12	0.516034485	+/- Scale factor * (S / \sqrt{n})
13	0.516082033	0.001063288
14	0.520432514	
15	0.519474935	
16	0.51189244	Lower bound
17	0.517458395	0.51678604
18	0.518170709	
19	0.517412674	Upper bound
20	0.516660272	0.518912617
21	0.521097394	
22	0.519237953	
23	0.516866164	
24	0.512226835	
25	0.520994809	
26	0.519302828	
27	0.518061165	
28	0.520910787	
29	0.515990025	
30	0.518503144	

7 Servers

Seed	Mean Response Time	Sample Mean
1	0.515411347	0.515322555
2	0.513182531	
3	0.516323247	Sample SD
4	0.515375042	0.001780174
5	0.514134054	
6	0.516335001	alpha
7	0.515388019	0.05
8	0.514500065	
9	0.512841075	scale factor (t 29, 0.975)
10	0.515036258	2.04523
11	0.513948175	
12	0.518645305	+/- Scale factor * (S / \sqrt{n})
13	0.513379182	0.000664728
14	0.513353974	
15	0.516150726	
16	0.515437287	Lower bound
17	0.514619214	0.514657826
18	0.515071589	
19	0.513762196	Upper bound
20	0.51538042	0.515987283
21	0.519487834	
22	0.516566629	
23	0.515849065	
24	0.51644742	
25	0.513766393	
26	0.512436924	
27	0.514437123	
28	0.515623028	
29	0.517458794	
30	0.519328719	

8 Servers

Seed	Mean Response Time	Sample Mean
1	0.517892523	0.521656586
2	0.5222454	
3	0.525498342	Sample SD
4	0.520281803	0.001783123
5	0.523288245	
6	0.522867144	alpha
7	0.522848173	0.05
8	0.52267093	
9	0.521174723	scale factor (t 29, 0.975)
10	0.523640597	2.04523
11	0.521642368	

12	0.519979925	+/- Scale factor * (S / \sqrt{n})
13	0.518126153	0.000665829
14	0.521842568	
15	0.522346787	
16	0.520702666	Lower bound
17	0.523664192	0.520990757
18	0.519905695	
19	0.521143364	Upper bound
20	0.524575655	0.522322416
21	0.521755262	
22	0.523228462	
23	0.520509697	
24	0.52192408	
25	0.5213105	
26	0.523209305	
27	0.520556322	
28	0.52199049	
29	0.519157887	
30	0.51971833	

9 Servers

Seed	Mean Response Time	Sample Mean
1	0.532533926	0.530814307
2	0.530896274	
3	0.528453199	Sample SD
4	0.532117399	0.001386869
5	0.530169895	
6	0.532422614	alpha
7	0.533467524	0.05
8	0.532582667	
9	0.530076181	scale factor (t 29, 0.975)
10	0.530547282	2.04523
11	0.531262269	
12	0.531956396	+/- Scale factor * (S / \sqrt{n})
13	0.531663306	0.000517866
14	0.527761558	
15	0.532351624	
16	0.531459757	Lower bound
17	0.531878382	0.530296442
18	0.528406667	
19	0.529181691	Upper bound
20	0.529580274	0.531332173
21	0.530920419	
22	0.528903073	
23	0.531481745	
24	0.529880904	

25	0.531507067
26	0.53049731
27	0.530368592
28	0.530209344
29	0.531223934
30	0.53066794

10 Servers

Seed	Mean Response Time	Sample Mean
1	0.539673373	0.541918933
2	0.543188946	
3	0.541158319	Sample SD
4	0.542584734	0.001630058
5	0.539650387	
6	0.540318834	alpha
7	0.541489529	0.05
8	0.540681794	
9	0.541800349	scale factor (t 29, 0.975)
10	0.543544824	2.04523
11	0.542381336	
12	0.544311077	+/- Scale factor * (S / \sqrt{n})
13	0.541246143	0.000608674
14	0.541673835	
15	0.54326916	
16	0.542493609	Lower bound
17	0.542638447	0.54131026
18	0.541528388	
19	0.543732403	Upper bound
20	0.538383764	0.542527607
21	0.543832035	
22	0.541411981	
23	0.544446093	
24	0.540692801	
25	0.540613651	
26	0.53954409	
27	0.544608901	
28	0.540431381	
29	0.543735935	
30	0.542501884	

We can see that there is minimal overall between adjacent number of servers so we do not need to increase n. If wish to confirm which is better out of 10 vs 5 servers we would need to increase our number of replications or increase our simulation time. However we can see the fastest number of servers with 95% confidence so we will stay with $n = 30$ and sim time set to 5000 - 15000

Comparing Response Times

From the upper and lower bounds calculated above and the charts below we can conclude with 95% confidence that the optimal number of servers is 7 with a 95% probability of mean response time between $[0.514658, 0.515987]$. We can conclude this because we have no overlap in the confidence intervals closest to the 7 server interval. Thus, this is the correct choice with $>95\%$ probability as there is space between 7 server interval and the next fastest interval (6 servers). This means we could increase the width of the interval and the confidence percentage above 95% while still avoiding any overlap between 6 and 7 intervals. In order from best to worst response time the order of servers is: 7, 6, 8, 9, 10*, 5*, 4, 3

*Note: that we cannot see any clear difference between 5 and 10 servers

